



Part 2

Processors

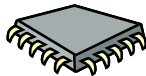


Computer Processors

The Brains of the Box

Computer Processors

- The *Central Processing Unit (CPU)* is the most complex part of a computer
- In fact, it is the computer
- It works far different from a high-level language



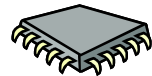
2/2/2017

Sacramento State - Cook - CS&35 - Spring 2017

3

Components of a Processor

- *Execution Unit (EU)*
 - performs calculations & logic
 - registers hold data
- *Control Logic Unit (CLU)*
 - reads and decodes instructions
 - talks to other components



2/2/2017

Sacramento State - Cook - CS&35 - Spring 2017

4

Execution Unit

- Contains the hardware that *executes* tasks (your programs)
- Different in many processors
- Modern processors often use multiple execution units to execute instructions in parallel to improve performance

2/2/2017

Sacramento State - Cook - CS&35 - Spring 2017

5

Execution Unit – The ALU

- The *Arithmetic Logic Unit* performs all calculations and comparisons
- Processor often contains special hardware for integer and floating point



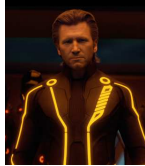
2/2/2017

Sacramento State - Cook - CS&35 - Spring 2017

6

Control Logic Unit (CLU)

- Controls the processor
- Determines when instructions can be executed
- Controls internal operations
 - fetch and execute each instruction
 - and store result of each instruction



2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

7

CLU Over Time

- In early processors...
 - CLU was a very small fraction of the hardware
 - EU and the registers took most of the space
- New processors...
 - complex control unit one of the more difficult parts of a processor to design
 - has increased in its percentage of the processor hardware

2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

8

Computer Processors

- Over time, thousands of processors were developed
- Examples:
 - Intel x86
 - IBM PowerPC
 - MOS 6502
 - ARM



2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

9

High-Level Programming

- You are used to writing in high level programming languages such as C#, Java, Visual Basic, etc...
- These are *third-generation languages* – and are designed to isolate you from architecture of the machine



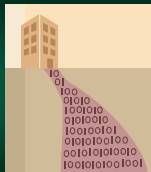
2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

10

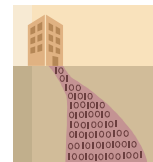
Registers

Where the work is done



Registers

- In high level languages, you put active data into variables
- However, it works quite different on processors
- All computations performed are done in *registers*



2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

12

What are registers used for?

- Registers are used to store anything the processor needs to keep track of
- Examples:
 - the result of calculations
 - status information
 - memory location of the running program
 - and much more...

2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

13

What – exactly – is a register?

- A register is a memory location located on the processor itself
- Think of it as a special "variable"
- Designed to be fast
- Some are accessible and usable by a programs, but many are hidden.

2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

14

General Purpose Registers

- *General Purpose Registers (GPR)* don't have a specific purpose
- They are designed to be used by programs – however they are needed
- Often, you must use registers to perform calculations

2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

15

Some Special Registers

- Program counter
 - controls the current memory location of the running program
 - privileged – only the processor and OS can change it
- Stack pointer
 - tracks the top of the system stack
 - you can modify this ... *with care...*

2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

16

Register Files

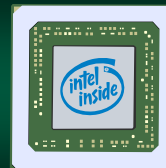


- All the related registers are grouped into a *register file*
- Different processors access and use their register files in very different ways
- Some processors support multiple files

2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

17

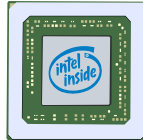


Original x86
Registers

It was simple at first...

Original x86 Registers

- First "x86" was the Intel 8086 released in 1978
- Attributes:
 - 16-bit processor (registers were 16-bit)
 - 16 registers
 - can access of 1MB of RAM



2/2/2017

Sacramento State - Cook - CS&IS - Spring 2017

19

Intel x86 Registers

- 8 Registers can be used by your programs
 - Four General Purpose: **AX, BX, CX, DX**
 - Four pointer index: **SI, DI, BP, SP**
- The remaining 8 are restricted
 - Six segment: CS, DS, ES, FS, GS, SS
 - One instruction pointer: IP
 - One status register – used in computations

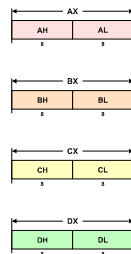
2/2/2017

Sacramento State - Cook - CS&IS - Spring 2017

20

Original General Purpose Registers

- However, back then (and now too) it is very useful to store 8-bit values
- So, Intel chopped 4 of the registers in half
- These registers have generic names of A, B, C, D



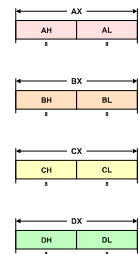
2/2/2017

Sacramento State - Cook - CS&IS - Spring 2017

21

Original General Purpose Registers

- The first and second byte can be used separately or used together
- Naming convention
 - high byte has the suffix "H"
 - low byte has the suffix "L"
 - for both bytes, the suffix is "X"



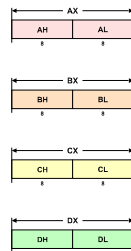
2/2/2017

Sacramento State - Cook - CS&IS - Spring 2017

22

Original General Purpose Registers

- This essentially doubled the number of registers
- So, there are:
 - four 16-bit registers or
 - eight 8-bit registers
 - (and any combination you can think off)

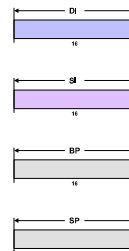


2/2/2017

Sacramento State - Cook - CS&IS - Spring 2017

23

Last the 4 Registers



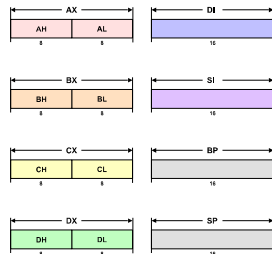
- The remaining 4 registers were not cut in half
- Used for storing indexes (for arrays) and pointers
- Their purpose
 - DI – destination index
 - SI – source index
 - BP – base pointer
 - SP – stack pointer

2/2/2017

Sacramento State - Cook - CS&IS - Spring 2017

24

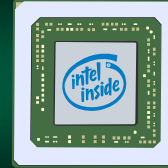
Original 16-Bit Registers



2/2/2017

Sacramento State - Cook - CS&IS - Spring 2017

25

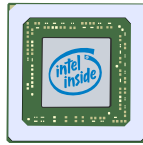


Evolution to 64 Bit Registers

This is going to hurt...

Original x86 Registers

- The x86 processor has evolved continuously over the last (nearly) 4 decades
- It jumped to 32-bit, and then finally 64-bit
- The result is many of the registers have strange names



2/2/2017

Sacramento State - Cook - CS&IS - Spring 2017

27

Evolution to 32-bit

- When the x86 moved into the 32-bit era, Intel expanded the registers to 32-bit
 - the 16-bit ones still exist
 - ... but also have a 32-bit version
 - they have the prefix *"e" for extended*
- New instructions were added (to use them)

2/2/2017

Sacramento State - Cook - CS&IS - Spring 2017

28

Evolution to 64-bit

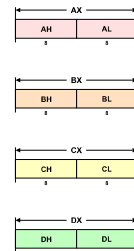
- Once again, the processor evolved – now to 64-bit
- The registers were extended again
 - the 64-bit have the prefix *"r" for register*
 - 8 additional registers were added
 - also, it is now possible to get 8-bit values from all registers (hardware is more consistent!)

2/2/2017

Sacramento State - Cook - CS&IS - Spring 2017

29

Original Registers

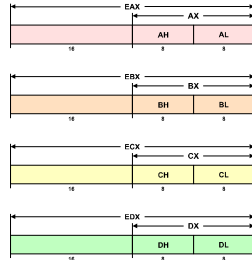


2/2/2017

Sacramento State - Cook - CS&IS - Spring 2017

30

Expansion to 32-bit

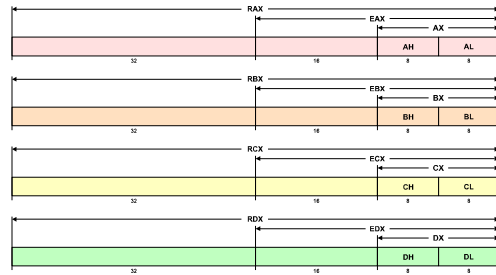


2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

31

Expansion to 64-bit

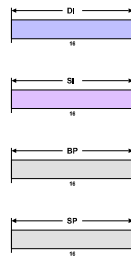


2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

32

Original Registers

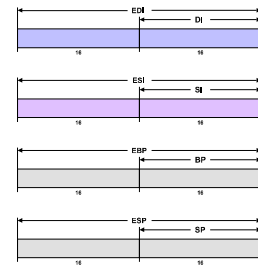


2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

33

Expansion to 32-bit

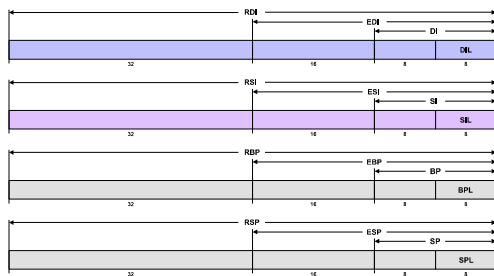


2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

34

Expansion to 64-bit

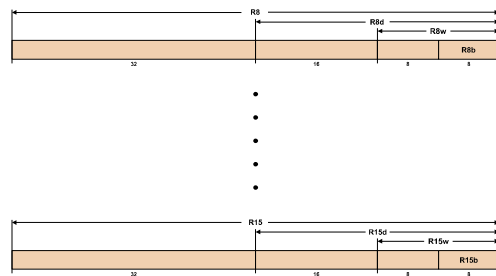


2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

35

New 64-bit Registers



2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

36

64-Bit Register Table

Register	32-bit	16-bit	8-bit High	8-bit Low
rax	eax	ax	ah	al
rbx	ebx	bx	bh	bl
rcx	ecx	cx	ch	cl
rdx	edx	dx	dh	dl
rsi	esi	si		sil
rdi	edi	di		dil
rbp	ebp	bp		bpl
rsp	esp	sp		bsl

2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

37

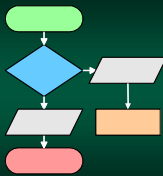
64-Bit Register Table

Register	32-bit	16-bit	8-bit High	8-bit Low
r8	r8d	r8w		r8b
r9	r9d	r9w		r9b
r10	r10d	r10w		r10b
r11	r11d	r11w		r11b
r12	r12d	r12w		r12b
r13	r13d	r13w		r13b
r14	r14d	r14w		r14b
r15	r15d	r15w		r15b

2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

38

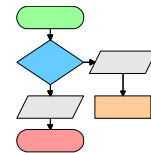


Instructions

Your programs are simple

Instructions

- Processors **do not** have the constructs you find in high-level languages
- Examples:
 - Blocks
 - If Statements
 - While Statements
 - ... etc



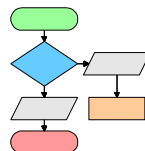
2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

40

Instructions

- Processors can only a series of simple tasks
- These are called *instructions*
- Examples:
 - add two values together
 - move a value
 - jump to a memory location



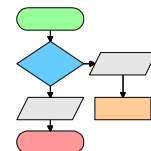
2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

41

Instructions

- These instructions are used to create **all** logic needed by a program
- We will cover how to do this during the semester

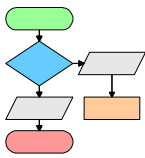


2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

42

Processor Instruction Set



- A processor's *instruction set* defines all the available instructions
- The instructions and their respective formats are very different for each processor

2/2/2017

Sacramento State - Cook - CS&35 - Spring 2017

43

What – exactly – is an instruction?

- An instruction is a series of bytes that contain everything the processor needs to know to do something
- An instruction must specify:
 - *operation* – what to do
 - *operands* – what data is to be used

2/2/2017

Sacramento State - Cook - CS&35 - Spring 2017

44

Operation Codes

- Each instruction has an *operation code (Opcode)*
- This a unique value that specifies the exact operation to be performed by the processor
- Assemblers use friendly names for called *mnemonics*

2/2/2017

Sacramento State - Cook - CS&35 - Spring 2017

45

Types of Operations (Opcodes)

- Data Transfer
- Program Flow Control
- Arithmetic and Logic operations
- Input and Output Instructions

2/2/2017

Sacramento State - Cook - CS&35 - Spring 2017

46

Operations: Data Transfer

- One of the most common tasks is moving data to and from registers
- Classified into three categories:
 - *loading* a register with data in memory
 - *storing* data in a register into memory
 - *transferring* data between registers

2/2/2017

Sacramento State - Cook - CS&35 - Spring 2017

47

Operations: Control Flow

- Processors do not support blocks, If Statements, etc...
- Instead, you must jump around code you don't want to execute
- This is the same of idea of GoTo Statements



2/2/2017

Sacramento State - Cook - CS&35 - Spring 2017

48

Operations: Arithmetic and Logic

- Many operations are used to modify data such as arithmetic, comparisons, and shifting
- Comparison is in this category
 - when two operands are compared – often one is **subtracted** from the other
 - result sets Boolean flags – *more on this later!*



2/2/2017

Sacramento State - Cook - CSIS 35 - Spring 2017

49

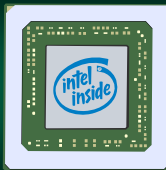
Operations: Input and Output

- There are also instructions that are designed to talk to ports, hard drives and other components
- However, in modern systems, these are **privileged** and only usable by the operating system
- You will use interrupts to tell the operating system to input/output data

2/2/2017

Sacramento State - Cook - CSIS 35 - Spring 2017

50

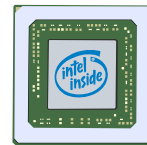


Basic Intel x86 Instructions

Feel the pow-wah of the x86!

Basic Intel x86 Instructions

- Each x86 instruction can have up to 2 operands
- Operands in x86 instructions are **very** versatile
- Often each argument can be either a memory address, register or an immediate value



2/2/2017

Sacramento State - Cook - CSIS 35 - Spring 2017

52

Types of Operands

- Registers
- Memory address
- Register pointing to memory
- A constant stored with the instruction – this is called an **immediate**

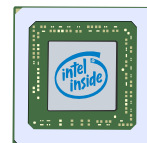
2/2/2017

Sacramento State - Cook - CSIS 35 - Spring 2017

53

Intel x86 Instruction Limits

- There are some limitations...
- Some instructions must use an immediate
- Some instructions require a **specific** register to perform calculations



2/2/2017

Sacramento State - Cook - CSIS 35 - Spring 2017

54

Intel x86 Instruction Limits

- A register must always be involved
 - processors use registers for all activity
 - both operands cannot access memory at the same time
 - *the processor has to have it at some point!*
- Also, obviously, the receiving field cannot be an immediate value

2/2/2017

Sacramento State - Cook - CS:35 - Spring 2017

55

Instruction: Move

- The x86 Move Instruction combines load, store, and register transfer logic
- It is one of the most common instructions used in programs (true of all processors)
- Remember how often you use the assignment statement in C / Java?

2/2/2017

Sacramento State - Cook - CS:35 - Spring 2017

56

Instruction: Move

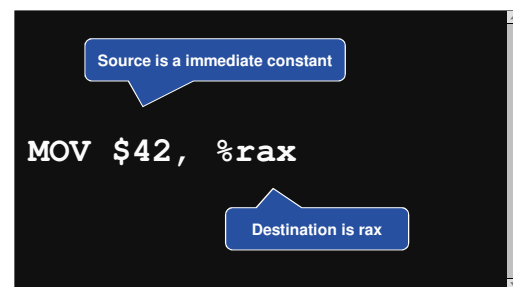


2/2/2017

Sacramento State - Cook - CS:35 - Spring 2017

57

Example: Move immediate



2/2/2017

Sacramento State - Cook - CS:35 - Spring 2017

58

Example: "A" Register

```
# So many options!

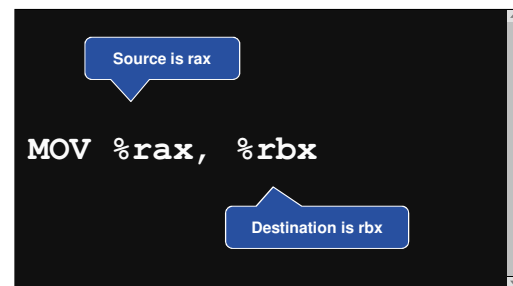
mov $42, %al    #low byte
mov $13, %ah    #high byte
mov $47, %ax    #both bytes
```

2/2/2017

Sacramento State - Cook - CS:35 - Spring 2017

59

Example: Move register to register



2/2/2017

Sacramento State - Cook - CS:35 - Spring 2017

60

Example: Move register to memory

Source is rax

```
MOV %rax, counter
```

Memory location named 'Counter'

2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

61

Instruction: Add & Subtract

- The Add and Subtract instructions take two operands and store the result in the second operand
- This is the same as the $+=$ and $-=$ operators used in Visual Basic .NET, C, C++, Java, etc...



2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

62

Instruction: Add

Immediate, Register, Memory

```
ADD value, target
```

Register, Memory

2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

63

Example: Move register to memory

Move memory into rax

```
MOV counter, %rax  
ADD $2, %rax
```

rax += 2

2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

64

Instruction: And & Or

- The Logical And and Logical Or instructions take two operands and stores the result in the second operand
- This is the same as the $\wedge=$ and $\vee=$ operators used in C, C++, Java, etc...



2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

65

Instruction: Logical And

```
AND value, target
```

2/2/2017

Sacramento State - Cook - CS& 35 - Spring 2017

66

Example: Logical Or

```
#Convert 5 to ASCII '5'  
MOV $5, %rax  
OR $0x30, %rax
```

0011 0000

2/2/2017

Sacramento State - Cook - CSC 35 - Spring 2017

67

Instruction: Call

- The Call Instruction transfers control to a memory location (a subroutine)
- Subroutines are analogous to the functions you wrote in Java
- Once it completes, execution continues after the call

2/2/2017

Sacramento State - Cook - CSC 35 - Spring 2017

68

Instruction: Call

CALL *address*

Usually a label – an constant that holds an address

2/2/2017

Sacramento State - Cook - CSC 35 - Spring 2017

69

Example: Print an integer

```
#This is using the CSC35 library
```

```
MOV $42, %rax  
CALL PrintInt
```

This name is an address

2/2/2017

Sacramento State - Cook - CSC 35 - Spring 2017

70