




## Part 6

### Memory & Addressing

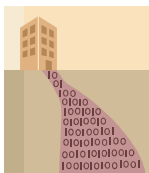


## What is Memory?

Its... um.... I forgot....

## Computer Memory

- Assembly offers you vast control over memory
- Understanding it...
  - is vital to becoming a great assembly programmer
  - and understanding computer architecture



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

3

## What is Memory?

- Memory is essentially a long list of bytes
- Memory is sometimes referred to as *storage*
- This is because it stores both running programs and their related data

Memory	
0	01000100
1	01000011
2	01101111
3	01101111
4	01101011

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

4

## Memory Addresses

- Memory is divided into a storage locations that can hold 1 byte (8 bits) of data
- Each byte has an *address*
  - unique value that refers to that specific byte
  - used to locate the exact byte the processor wants

Memory	
0	01000100
1	01000011
2	01101111
3	01101111
4	01101011

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

5

## Metaphor for Memory

- Think of memory as a *set of mailboxes*
- Each mailbox can contain a piece of data (byte)
- Each mailbox has a unique number



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

6

## Metaphor for Memory

- ... or think of memory as a *group of boxes*
- Each box belongs to the same variable
- Each box has a unique number



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

7

## Endianness

The "proper" order of things

## So Many Bytes...

- On a 64-bit system, each word consists of 8 bytes
- So, when any 64-bit value is stored in memory, each of those 8 bytes must be stored
- However, question remains: *What order do we store them?*



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

9

## So Many Bytes...

- Do we store the least-significant byte (LSB) first, or the most-significant (MSB)?
- As long as a system always follows the same format, then there are no problems
- ... but different systems use different approaches

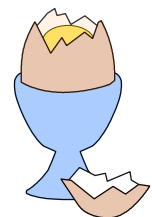
3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

10

## Big Endian vs. Little Endian

- Big-Endian approach
  - store the MSB first
  - used by Motorola & PowerPC
- Little-Endian approach
  - store the LSB first
  - used by Intel
  - appears "backwards" in editors



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

11

## Example Unsigned Integer (4 Byte)

3,721,182,122

DD	CC	BB	AA
----	----	----	----

Most significant Byte

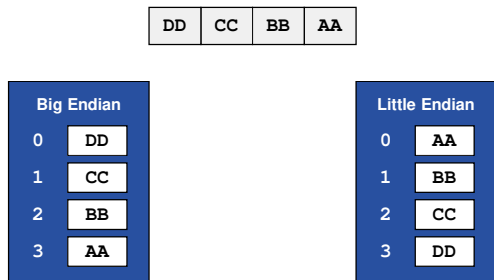
Least significant Byte

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

12

## Big Endian vs. Little Endian



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

13

## No "End" to Problems

- *There is a problem...*  
if two systems use different formats, data will be interpreted incorrectly!
- For example:
  - a **little**-endian system reads a value stored in **big**-endian
  - a **big**-endian system reads a value stored in **little**-endian



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

14

## No "End" to Problems

- So, whenever data is read from secondary storage, you cannot assume it will be in your processor's format
- This is compounded by file formats (gif, jpeg, mp3, etc...) which are also inconsistent



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

15

## Endianness in File Formats

- Adobe Photoshop - Big Endian
- Windows Bitmap (.BMP) - Little Endian
- GIF - Little Endian
- JPEG - Big Endian


3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

16

## So... who is correct?

- Is the Intel x86 (little endian) or the PowerPC (big endian) correct?
- In reality...
  - neither side is correct
  - both formats are equally correct



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

17


## Gulliver's Travels



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

18




## Addressing Modes

How to interact with memory

## Addressing Modes


- Processor instructions often need to access memory to read values and store results
- So far, we have used registers to read and store single values
- However, we need to:
  - access items in an array
  - follow pointers
  - and more!



3/13/2017 Sacramento State - Cook - CSc 35 - Spring 2017 20

## Addressing Modes

- How a processor can locate and read data from memory is called an *addressing mode*
- Information combined from registers, immediates, etc... to create a target address
- Modes vary greatly between processors



3/13/2017 Sacramento State - Cook - CSc 35 - Spring 2017 21

## 4 Basic Addressing Modes

- Value stored in a register
- Memory address specified in the instruction
- Memory address pointed to by a register
- Immediate (part of instruction after the opcode bits)

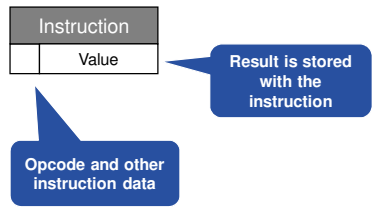
3/13/2017 Sacramento State - Cook - CSc 35 - Spring 2017 22

## Immediate Addressing

- Immediate addressing is one of the most basic modes found on a processor
- Often a value is stored as part of the instruction
- As the result, it is *immediately* available
- Very common for assigning constants

3/13/2017 Sacramento State - Cook - CSc 35 - Spring 2017 23


## Immediate Addressing



3/13/2017 Sacramento State - Cook - CSc 35 - Spring 2017 24

## Register Addressing

- Register addressing is used in practically all computer instructions
- A value is read from or stored into one of the processor's registers
- Instruction contains the register's **number**




3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

25

## Register Addressing



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

26

## Direct Addressing


- In **direct** addressing, the processor reads data directly from the computed address
- Commonly used to:
  - get a value from a "variable"
  - read items in an array
  - etc...

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

27

## Direct Addressing

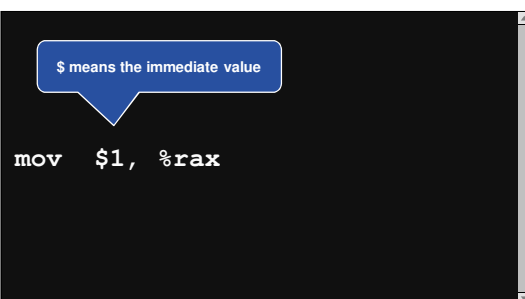


3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

28

## Example: Immediate

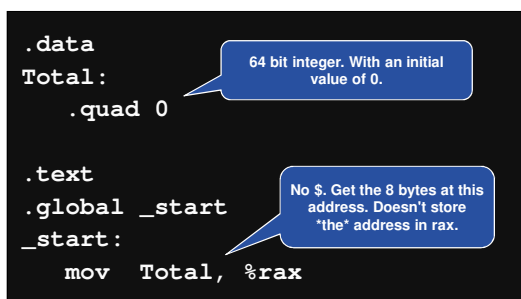


3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

29

## Example: Direct



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

30

## Register Indirect Addressing


- *Register Indirect* uses a register is used to store the address
- Same concept as a *pointer*
- Because the address is in a register...
  - processor does have to go to memory get it
  - it is just as fast as direct addressing
  - ... and very common

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

31

## Register Indirect Addressing



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

32

## Relative Access

- In *relative access*, a value is added to a system register (e.g. program counter)
- Advantages:
  - instruction can just store the *difference* (in bytes) from the current instruction address
  - takes less storage than a full 64-bit address
  - it allows a program to be stored anywhere in memory – *and it will still work!*

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

33

## Relative Addressing

- Often used in conditional jump statements
  - only need the to store the number of bytes to jump – either up or down
  - so, the instruction only stores the value to add to the program counter
  - practically all processors us this approach
- Also used to access local data – load/store

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

34



## Behind the Scenes of Arrays

All the mystery is revealed!

## Arrays

- Computers do not have an 'array' data type
- So, how do you have array variables?
- When you create an array...
  - you allocate a block of memory
  - each element (cell) is located sequentially in memory – one right after each other



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

36

## Arrays

- Every byte in memory has an address
  - ... as does every element in an array
  - to get an array cell, we merely need to *compute* the address
- The "index" and "scale" addressing features are designed for arrays
- ... well, that and *any* block of memory

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

37

## Array Math Example

- Start of our block of memory (buffer) is at address **2000**
- The first array cell is at 2000
- Arrays consists of bytes...
  - the second is at **2001**
  - the third is at **2002**
  - the fourth at **2003**
  - etc...

2000	H
2001	e
2002	l
2003	l
2004	o

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

38

## Array Math Example – 32 bit

- However, what if we are storing 32-bit integers?
- A 32-bit integer takes **4** bytes in memory
- So, as a result, *each cell will require 4 bytes of memory*

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

39

## Array Math Example – 32 bit

- First cell uses 2000... 2003
- Since each cell is 4 bytes...
  - the second is at **2004**
  - the third is at **2008**
  - the fourth at **2012**
  - etc...

2000	F0A3
2004	042B
2008	C1F1
2012	0D0B
2016	9C2A

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

40

## Array Math Example – 64 bit

- The case with 64-bit integers is exactly the same
- A 64-bit integer takes **8** bytes in memory
- So, as a result, *each cell will require 8 bytes of memory*

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

41

## Array Math Example – 64 bit

- First cell uses 2000... 2007
- Since each cell is 8 bytes...
  - the second is at **2008**
  - the third is at **20016**
  - the fourth at **2024**
  - etc...

2000	F0A3
2008	042B
2016	C1F1
2024	0D0B
2032	9C2A

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

42

## Behind the Scenes...

- So, when an array element is read, internally, a mathematical equation is used
- It takes into account the start of the first cell, the array index, and the size of each element

`Start of Buffer + (Index * Size)`

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

43

## Behind the Scenes...

- *This is why the C Programming Languages uses zero as the first array element*
- If zero is used with this formula, it gets the start of the buffer

`Start of Buffer + (Index * Size)`

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

44

## Behind the Scenes...

- Java uses zero-indexing because C does
- ... and C does so it can create efficient assembly!

`Start of Buffer + (Index * Size)`

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

45



## Addressing on the x86

Grabbing any byte

## Addressing on the x86

- The Intel x86 supports direct, indirect, indexing and scaling
- So, the Intel is very versatile in how it can access memory
- This is typical of CISC-ish architectures



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

47

## Effective Addresses

- Using the addresses stored in memory, registers, etc... is useful in programs
- Often programs contain *groups* of data
  - fields in an abstract data type
  - cells in an array
  - entries in a large table etc...



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

48



## Effective Addresses

- Processors have the ability to create an *effective address* by combining data
- How it works:
  - starts with a base address
  - then adds a value (or values)
  - finally, uses this temporary value as the actual address



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

49

## Terminology

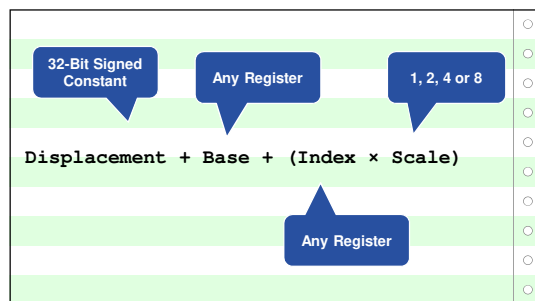
- Base-address* is the initial address
- Displacement (aka offset)* is a constant (immediate) that is added to the address
- Index* is a *register* added to the address
- Scale* used to multiply the index before adding it to the address

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

50

## x86 Effective Address Formula




3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

51

## Behind the Scenes...

- But wait, doesn't that formula look familiar?
- The addressing term "scale" is basically equivalent to "size" in this example
- Addressing helps us use arrays!



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

52

## Addressing Notation in Assembly

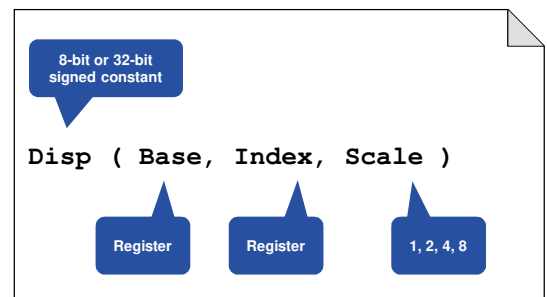
- The AT&T / GAS notation allows you to specify the full addressing
- The notation is a tad terse, and the alternative, Intel notation, is easier to read
- However...
  - you will get used to it quite quickly
  - look at what you can read already!

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

53

## AT&T / GAS Operand Notation



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

54

## AT&T / GAS Notation

Mode	Syntax	Example
Direct	<b>Address</b>	<code>mov <b>address</b>, %rdx</code>
Direct Indexed	<b>Address (Index)</b>	<code>mov <b>address(%rax)</b>, %rdx</code>
Register Indirect	<b>(Register)</b>	<code>mov <b>(%rax)</b>, %rdx</code>
Register Indirect Indexed	<b>(Register, Index)</b>	<code>mov <b>(%rax, %rbx)</b>, %rdx</code>

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

55

## Addressing Notation in Assembly

- When you write an assembly instruction...
  - you specify all 4 four addressing features
  - however, notation fills in the "missing" items
- For example: for direct addressing...
  - Displacement → Address of the data
  - Base → Not used
  - Index → Not used
  - Scale → 1, which is irrelevant without an Index

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

56

## How Many Bytes

- When you store data into a register, the assembler knows (*by looking at the size of the register*) how much is going to be accessed
- However, when using addressing,
  - it sometimes is not obvious if you are accessing a byte, 2 bytes, etc...
  - this will cause a very cryptic error

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

57

## How Many Bytes

- To address this issue, AT&T/GAS notation places a single character after the instruction name
- This suffix will tell the assembler how many bytes will be accessed during the operation

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

58

## How Many Bytes

Suffix	Meaning
<b>b</b>	byte
<b>s</b>	short (2 bytes)
<b>l</b>	long (4 bytes)
<b>q</b>	quad (8 bytes)

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

59

## Example: Direct Index

Using the EDI register for indexing, but you can use any GP register

```
mov $1, %rdi
movb $33, Text(%rdi)
```

ASCII 33 → !

Offset	Value
0	H
1	!
2	L
3	L
4	O

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

60

### Example: Direct Index (Scale of 2)

```
mov $1, %rdi
movb $33, Text(,%rdi,2)
```

Each "cell" is 2 bytes

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

61

### Example: Direct Index (Scale of 4)

```
mov $1, %rdi
movb $33, Text(,%rdi,4)
```

Each "cell" is 4 bytes

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

62

### Example: Register Indirect

```
mov $Text, %rax
movb $33, (%rax)
```

The value of Text - an address

Indirect. Base is rax

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

63

### Example: Register Indirect Index

```
mov $Text, %rax
mov $1, %rdi
movb $33, (%rax,%rdi)
```

Base

Index

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

64

### Example: Reg Indirect Index (Scale 2)

```
mov $Text, %rax
mov $1, %rdi
movb $33, (%rax,%rdi,2)
```

Scale

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

65

### Example: Reg Indirect Index (Scale 4)

```
mov $Text, %rax
mov $1, %rdi
movb $33, (%rax,%rdi,4)
```

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

66

## For Loop: 0 to 4

```

mov $0, %rdi

Loop:
    cmp $4, %rdi
    jg End

    movb $33, Text(%rdi)
    add $1, %rdi
    jmp Loop

End:

```

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

67

## For Loop: 0 to 4

```

mov $0, %rdi

Loop:
    cmp $4, %rdi
    jg End

    movb $33, Text(%rdi)
    add $1, %rdi
    jmp Loop

End:


```

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

68


## x86 Register Mastery



Choosing the right register (and mode)

## x86 Register Mastery

- x86 has 8-bit, 16-bit, 32-bit, and 64-bit registers
- They are different parts of the same register (e.g. ah, al, ax, rax are the "A" register)
- Using the correct one is **vital** to making your program work



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

70

## x86 Register Mastery


- When you load/store data, the register will grab as many bytes as it can store
- So...
  - 8-bit register will access 1 byte
  - 16-bit register will access 2 bytes
  - 32-bit register will access 4 bytes
  - 64-bit register will access 8 bytes
- Using the wrong size can cause problems

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

71

## Expansion to 64-bit




3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

72

## Expansion to 64-bit




3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

73

## New 64-bit Registers



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

74

## Example Program

```
.data
Message:
    .ascii "Hello"

.text
.global _start

_start:
    mov Message, %eax
```

Creates 5 bytes to store Hello

eax is 32-bit (4 bytes)

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

75

## Example Program

```
.data
Message:
    .ascii "Hello"

.text
.global _start

_start:
    mov Message, %eax
```

Message: 48 H, 65 e, 6C l, 6C l, 6F o

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

76

## Example Program

```
.data
Message:
    .ascii "Hello"

.text
.global _start

_start:
    mov Message, %eax
```

Message: 48 H, 65 e, 6C l, 6C l, 6F o

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

77

## Example Program

- In that example, we used a 32-bit register (eax) to read from the address "Message".
- It grabbed 4 bytes!
- If we wanted to compare a single character to another using 32-bit registers...
  - it would fail – we grabbed too much!
  - it would also compare those extra characters

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

78

## Example Program

```

.data
Message:
    .ascii "Hello"

.text
.global _start

_start:
    mov Message, %al
    
```



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

79

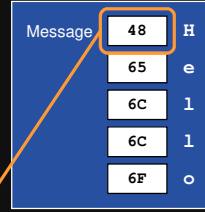
## This works, but gives a warning...

```

.data
Message:
    .ascii "Hello"

.text
.global _start

_start:
    movb Message, %rax
    
```



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

80



## Buffer Overflow

With Great Power  
Comes Great Responsibility

## Buffer Overflow

- Operating systems protect programs from having their memory / code damaged by another program
- However...operating systems don't protect programs from damaging *themselves*



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

82

## Buffers


- In memory, a running program's data is often stored next to its instructions
- Blocks of memory called *buffers* can store data (which can vary in size)
- Examples:
  - people's names
  - list of pet names
  - bytes in an image

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

83

## Buffer Overflow – How it Works



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

84

## Buffer Overflow



- It is possible to store too much information – resulting in a *buffer overflow*
- The extra bytes will overwrite part of the running program – changing it!

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

85

## Buffer Overflow – How it Works

Computer Memory




3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

86


## Tables



How to Organize Data

## Data Tables

- In assembly, you have full control of memory
- You can take advantage of these to create tables
- They can contain any data – from integers, to characters, to addresses



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

88

## ASCII Directive Creates a Table

Text:

```
.ascii "Hello"
```

Creates 5 bytes to store Hello. They are stored consequently

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

89

## Same Thing!

Text:

```
.byte 'H'  
.byte 'e'  
.byte 'l'  
.byte 'l'  
.byte 'o'
```

Created byte by byte

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

90

## Accessing Each Cell

Use register to hold table index

```
mov $1, %rdi
movb Text(%rdi), %ah
```

Text		
H	0	
E	1	
L	2	
L	3	
O	4	


3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

91

## Tables of Integers

- Tables can contain anything!
- Often, they are used to store integers & addresses (8 bytes on a 64-bit system)
- Just make sure to use the scale feature!



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

92

## Table of Long Integers

Values:

```
.quad 45
.quad 35
.quad 100
.quad 25
.quad 75
```

4 Bytes each

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

93

## Table of Long Integers

Values:

```
.quad 45
.quad 35
.quad 100
.quad 25
.quad 75
```

Values	45	0
	35	8
	100	16
	25	24
	75	32

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

94

## Accessing Each Cell

Table index 1

```
mov $1, %rdi
movl Values(,%rdi,8), %rax
```

Note the scale!

Values		
45	0	
35	8	
100	16	
25	24	
75	32	


3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

95

## Jump Table

- You can also jump to a value stored in a register
- ... which means you can create a jump table!



3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

96



## Table of Addresses

```
JumpTable:  
    .quad ReadInt  
    .quad PrintInt
```

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

97

## Calling a Register (a tad odd)

```
mov    $1, %rdi  
movl   JumpTable(,%rdi,8), %rbx  
  
call   *%rbx
```

AT&T notation  
requires an  
asterisk

3/13/2017

Sacramento State - Cook - CSc 35 - Spring 2017

98