




Part 10

Hardware




Instruction Execution

The Steps Involved

Instruction Execution

- When an instruction executes on a processor, a number of different tasks take place
- Typically, these tasks can be broken into 5 different steps



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

3

Five Steps in Execution

1. **Fetch** an instruction from memory
2. **Decode** it to determine what it is
3. **Read** the inputs from registers / memory
4. **Execute** for computations for instruction
5. **Write** the result into the registers / memory

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

4

1. Fetch the Instruction

- First, the processor fetches the instruction from the memory
- The result is stored in the Instruction Register
- Formally known as *Instruction Fetch (IF)*

Fetch

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

5

2. Decode the Instruction

- Second, the instruction is *decoded* to determine what it is and its operands
- Signals are sent to the execution unit as input
- Formally known as *Instruction Decode (ID)*

Decode

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

6

3. Read Inputs

- The Execution Unit then reads the values of the instruction
- These can be located in the instruction itself (immediate), register file, and from memory
- Formally known as *Memory Access (Mem)*

Read

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

7

4. Execute the Instruction

- Decoded instruction and input values are passed to the ALU
- Depending on the complexity of the instruction, some computations require multiple clock cycles
- e.g. multiplication requires more cycles and an add
- Formally: *Execute (EX)*

Execute

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

8

5. Write Result

- Finally, the result is written back into the register / memory
- Processor also updates flags and other state information such as the Program Counter
- Formally known as *Write Back (WB)*

Write

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

9

Instruction Execution




5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

10


Execution of 2 Instructions



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

11




Pipelining

Do multiple things at once

Pipelining

- *Pipelining* is an technique where multiple instructions are executed at the same time
- This greatly improves the speed (and efficiency) of a system



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

13

Pipelining

- It is invisible to the programmer... implemented by the hardware
- Pipelining is different from multi-core processors!
- Pipeline happens on one core



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

14

The Laundry Metaphor

- To understand the concept, typically a "laundry metaphor" is used
- Based on how we all do our laundry:
 - put the clothes in the washing machine
 - put the wet clothes in the dryer
 - fold the clothes and put them away

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

15

The Laundry Metaphor



Step 1. Put the clothes in the washer



Step 2. Put the wet clothes in the dryer




Step 3. Fold the clothes and put them away

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

16

Let's Do Our Laundry



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

17

Looking At Our Task

- It took us until 10pm to finish our laundry!
- But... some of our equipment was idle
 - when we were washing laundry, the dryer and folding counter was idle
 - when the dryer was being used, the washer and folding counter was idle
 - when the folding counter was being used, the washer and dryer were idle

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

18

The Pipelined Approach


- Let's overlap these loads of laundry
- We don't have to wait until the first load is done before we start the next one
- Better approach
 - put the Load 1 in the washing machine
 - after Load 1 is washed, we place it in the dryer
 - washer is now available... start load 2

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

19

Pipelined: Fold & Wash Together




5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

20

Fully Pipelined




5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

21

All Three Components in Use



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

22

So, is it faster?

- Note:** the actual time required to finish a load of laundry has *not* changed
- It still requires the same 3 steps and each takes one hour
- The speedup occurs because the different loads are executed in *parallel*

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

23

Analyzing Speedup

- So, how have we improved efficiency?
- In pipelining, the speedup is directly proportional to the number of steps (that we can do in parallel)

$$\text{Speedup}_{\text{pipelined}} = \text{Number of Steps}$$

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

24


Start-up and Wind-down

- At the beginning of the work load, the pipeline wasn't completely full
- So, the pipeline has to fill before we get our full speedup
- As the number of Loads increases...
 - pipeline is full for a larger fraction of the time
 - so the start-up and wind-down (finishing the last load) becomes meaningless

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

25




Pipelined Instructions

Modern Processors

Pipelined Instructions

- Just like the laundry metaphor, processors have different components that can be used at the same time
- On modern processors, practically all the hardware (all those transistors) are in continuous use



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

27

Pipelined Instructions

- Ideally, *nothing is ever idle!*
- Different stages of execution are pipelined
 - fetch
 - decode
 - read
 - execution
 - write back

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

28

Unpipelined Instructions

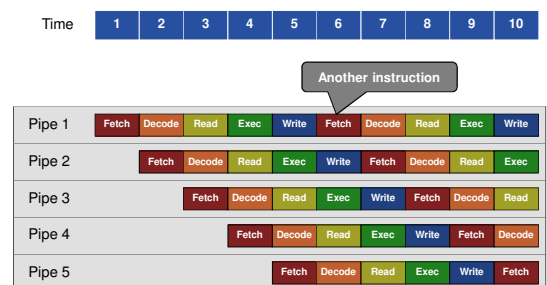


5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

29

5 Pipelines



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

30

Problems Arise

- The problem arises that different stages take different amounts of time
- This is not being measured in clock cycles – but in nanoseconds
- So, pipelining requires a mechanism to keep everything in sync


5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

31

Latches

- To implement pipelining, a execution is divided into stages
- These often exist, physically, on different sections of the processor




5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

32

Latches

- Designer places *latches* (aka *buffers*) between each section
- Instruction will not advance onto the next stage until the latch is lifted




5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

33

How Latches Work

- The latches read their inputs at the start of each cycle
- Store the data to be passed from them - remain constant throughout the rest of the cycle




5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

34

How Latches Work

- When a stage is ready...
 - it signals the prior stage
 - prior state sends the data
 - and becomes available




5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

35

Latched Instructions



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

36

Latched Instructions



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

37


Pipeline Read / Write Hazards



Call the plumber!

Pipeline Issues

- When you execute instructions in parallel, there are inherent dangers that must be handled
- Often, one instruction can interfere with another and cause strange side effects



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

39

Pipeline Issues

- There are 3 read/write dangers in pipelining
 - Read-After-Write (RAW)
 - Write-After-Write (WAW)
 - Write-After-Read (WAR)
- Each involves the writing of data to memory and / or registers
- These are known as *data dependencies*


5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

40

The Result of the Hazards

- These hazards will always result in the program malfunctioning – *even though it might be written correctly*
- The hardware fails... not the software



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

41

Read-After-Write (RAW)

- A *Read-After-Write Hazard* is caused when one instruction reads data before the previous one has written it
- Instruction **B** reads an operand before Instruction **A** writes it



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

42

Read-After-Write (RAW)

- The result is that the instruction **B** has read old data – before **A** changed it
- It is though instruction **A** doesn't exist!



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

43

Read-After-Write (RAW) Example

```
MOV $12, %rax
```

```
ADD $5, %rax
```

```
SUB %rax, value #Read output of add
```

If Sub reads rax before Add writes, it will subtract 12 rather than 17

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

44

Write-After-Read (WAR)

- A *Write-After-Read Hazard* is caused when one instruction writes data before the previous one has read it
- Instruction **B** writes an operand before Instruction **A** reads from it



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

45

Write-After-Read (WAR)

- The first instruction gets the wrong operand
- It should contain the value before Instruction **A** executes
- Instead, contains the value after instruction **B** executes



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

46

Write-After-Read (WAR) Example

```
MOV $6, %rax
```

```
SUB %rax, value  
ADD $1, %rax
```

If Add writes rax before Sub reads, it will subtract 7 rather than 6

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

47

Write-After-Write (WAW)

- A *Write-After-Write Hazard* is caused two instructions attempt to write data, and the later one writes it first
- Instruction **A** and **B** attempt to write the same operand, but **A** writes it last



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

48

Write-After-Write

- The operand contains the wrong result
- Should contain the value of instruction **B** but it contains the result of Instruction **A**
- It is though **B** never executed!



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

49

Write-After-Write (WAW) Example

```
MOV $42, %rax
MOV $31, %rax
```

If first MOV writes after the second MOV, rax will contain 42 rather than 31

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

50



Score Boarding

Solution to Data Dependency

Score Boarding

- *Score boarding* is a solution for register read stage structure hazards
- Processors track:
 - which registers will be written to by instructions in the pipeline
 - subsequent instructions determine whether their input registers are available




5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

52

Score Boarding Bits

- In score boarding, a "*presence*" bit (aka p-bit), is added to each register
- The presence bit records:
 - if the register is *full* – i.e. it is available for reading
 - or if the register is *empty* – i.e. waiting for an instruction to write its output



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

53

P-Bit Score Boarding



- P = 1, when data available for read
- After an instruction is going to write a register, P is set to 0 until it completes write-back

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

54

Operation: Input is Full (Ready)

- When an instruction enters the Read stage, the hardware checks to see if all of its input registers are *full (i.e. ready)*
- If so, the hardware reads the values of all the input registers

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

55

Operation: Input is Full (Ready)

- The hardware then marks *all* the *output* registers as *empty (busy)* – so no other instruction can read it until they are written
- The instruction then proceeds to the execute stage on the next cycle

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

56

Operation: Input is Empty

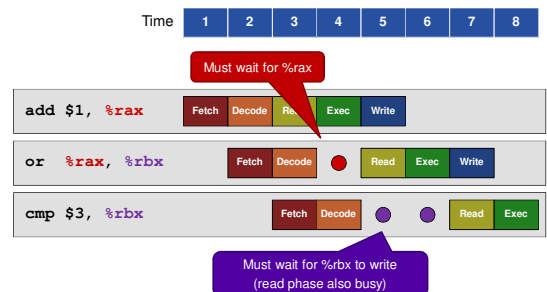
- If, in the Read Stage, the registers are *empty (busy)*, the instruction must *wait*
- The hardware holds the instruction, in the register read stage, until its input values become full
- In the meantime it *inserts bubbles* into the execute stage – it basically *stalls* it

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

57

Bubbles



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

58



Bypass Forwarding

Heads Up, the Data is Being Thrown!

Bypass Forwarding

- Another approach to handling data hazards is to, in some cases, skip the register file
- The result is *forwarded* to the next register's read stage



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

60

Bypass Forwarding

- Bypassing the write-back stage, we can eliminate data hazards
- Essentially,
 - we pass the ALU output from the first instruction directly to the second instruction
 - we skip the register file




5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

61

Paths for Result Forwarding




5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

62


Pipelined Instructions



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

63




Branch Hazard

Jumps and Pipelines are not friends

Branch Hazard

- So far, we covered register related hazards
- These are, unfortunately, not the only hazards that exist in pipelining
- Conditional branches cause problems




5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

65

Branch Hazard

- For example, what happens if several instructions are in the pipeline
- One instruction contains a conditional jump
- Do we just execute the instructions after it? They may never run!

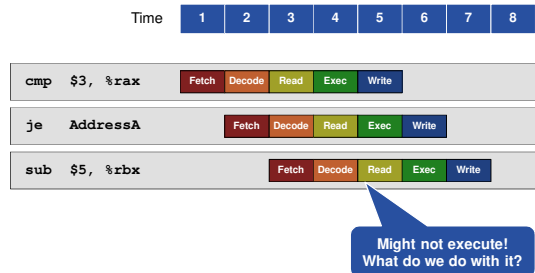


5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

66

Branch Hazard



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

67

Conditional Branches

- Unconditional and conditional branches, create data dependencies
- These are between...
 - branch instruction and the instruction fetch stage of the pipeline
 - branch instruction computes the address of the next instruction - that the instruction fetch stage should fetch

5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

68

Flushing

- After the conditional branch instruction is analyzed, the hardware changes the program counter
- All the instructions in the pipeline are invalid (since they will not execute)
- These pipelines are flushed – data discarded and restarted



5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

69

Flushing

- Naturally, this causes a huge delay for the new pipeline's startup latency
- So, *conditional branch penalty* has a huge effect on processor performance

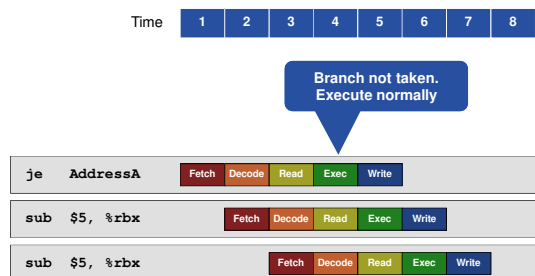


5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

70

Conditional Jump Not Taken

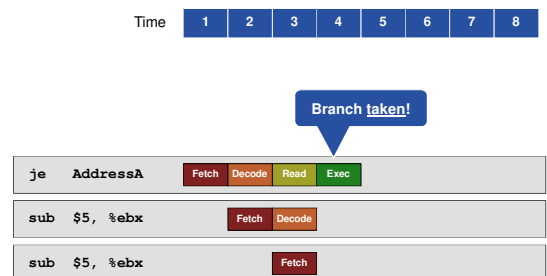


5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

71

Flush Required – 1 of 3




5/11/2017

Sacramento State - Cook - CSc 35 - Spring 2017

72

Flush Required – 2 of 3

Time 1 2 3 4 5 6 7 8



5/11/2017


Sacramento State - Cook - CSIS 35 - Spring 2017

73

Flush Required – 3 of 3

Time 1 2 3 4 5 6 7 8

Conditional branch penalty



5/11/2017

Sacramento State - Cook - CSIS 35 - Spring 2017

74